

Comparative Analysis of State Representations for Deep Q-Learning in the Snake Game

Sofie Amalie Roer Thorlund
Viktor Mekis Bach

January 2, 2026

Abstract

This paper investigates how state representation choice affects Deep Q-Network (DQN) performance in the Snake game. We compare two approaches: (1) a feature-engineered 10-dimensional vector capturing spatial relationships, and (2) a grid-based $30 \times 30 \times 3$ tensor processed via convolutional neural networks. Both agents were trained for 5,000 episodes on a 30×30 grid. Results demonstrate that the feature-based agent achieves superhuman performance (mean score 18.5, $7.4 \times$ human baseline) while the CNN agent shows limited learning (mean score 0.05). We attribute feature agent success to efficient encoding of task-relevant information and sample-efficient learning. CNN underperformance stems from shallow architecture, insufficient receptive field for global context, and inadequate training convergence. Our findings suggest that engineered features significantly outperform raw visual inputs for small-scale, fully-observable domains when training budgets are limited, challenging the assumption that CNNs universally excel at spatial reasoning tasks.

1 Introduction

Deep Reinforcement Learning (DRL) has achieved remarkable success in game-playing domains, from Atari games to the game of Go [1]. A critical yet often under-explored design choice in DRL is state representation: should agents learn from carefully engineered features that encode domain knowledge, or from raw sensory input such as images or grid representations? This fundamental question has significant implications for sample efficiency, generalization capability, and practical deployment.

The Snake game provides an ideal testbed for investigating this question. Unlike visually complex environments like Atari, Snake offers a compact, fully observable state space that eliminates confounding factors such as partial observability and perceptual noise. This controlled setting allows us to isolate the impact of representation choice on learning dynamics and final performance. Feature-based representations can efficiently encode critical information like snake position, food location, and imminent collision threats. In contrast, CNN-based approaches must discover these spatial patterns through hierarchical convolutions—a more

general approach that potentially enables transfer to previously unseen environments.

Existing literature demonstrates that CNNs excel in high-dimensional visual domains where hand-crafting features is infeasible [1, 5]. However, for games with clear, interpretable structure, engineered features often outperform raw inputs, particularly under constrained training budgets [9]. Our work contributes empirical evidence to this ongoing debate by systematically comparing both approaches in a controlled environment.

2 Research Questions

This paper investigates three interconnected research questions to understand how state representation affects deep reinforcement learning performance:

RQ1 - State Representation Impact: How do feature-based vs. CNN-based state representations affect DQN performance and training efficiency? We hypothesize that the compact, structured nature of Snake may favor feature representations despite CNNs’ theoretical advantages in spatial reasoning.

RQ2 - Human Baseline Comparison: How does DQN performance compare to human baseline? Establishing human performance provides crucial context for evaluating whether learned policies achieve practical utility and superhuman capability.

RQ3 - Strategy Emergence: What qualitative strategies emerge at different training stages between representations? Understanding behavioral evolution reveals how different state representations influence policy development and decision-making patterns.

3 Methods

3.1 Environment and Baseline

We implemented a Snake environment with a 30×30 grid. The agent controls a snake to eat randomly-placed food while avoiding walls and self-collision. Episodes terminate on collision or after 1,000 steps. The reward structure is: +10 for eating food, -10 for collision, -0.01 per step (encouraging efficiency).

3.2 Feature-Based Representation (DQN-Feature)

The feature state vector consists of 10 carefully designed dimensions capturing task-relevant information:

1. Relative food X-position (normalized, $\Delta x \in [-2, 2]$)
2. Relative food Y-position (normalized, $\Delta y \in [-2, 2]$)
3. Collision threat left (binary)
4. Collision threat front (binary)
5. Collision threat right (binary)
6. One-hot direction encoding: up (binary)
7. One-hot direction encoding: down (binary)
8. One-hot direction encoding: left (binary)
9. One-hot direction encoding: right (binary)
10. Normalized snake body length

These features encode minimum information necessary for optimal play: food location, immediate danger detection, current orientation, and progress tracking [6]. The vector is processed by a feed-forward network with 2 hidden layers (128 units, ReLU), totaling $\sim 18K$ parameters.

3.3 CNN-Based Representation (DQN-CNN)

The grid state is represented as a 3-channel 30×30 tensor, mimicking visual input:

- Channel 0: Snake body (1.0 for body segments, 0.0 otherwise)
- Channel 1: Snake head (1.0 for head position only)
- Channel 2: Food position (1.0 for food location)

This representation forces the agent to learn spatial patterns from raw grid data [1]. The CNN architecture processes this input through three convolutional layers with progressively increasing feature maps:

Layer sequence: Conv(3 $\rightarrow$ 32, k=5, p=2) $\rightarrow$ ReLU $\rightarrow$ MaxPool(2) [30 $\rightarrow$ 15]; Conv(32 $\rightarrow$ 64, k=3, p=1) $\rightarrow$ ReLU $\rightarrow$ MaxPool(2) [15 $\rightarrow$ 7]; Conv(64 $\rightarrow$ 128, k=3, p=1) $\rightarrow$ ReLU; Flatten $\rightarrow$ FC(256) $\rightarrow$ ReLU $\rightarrow$ Dropout(0.2) $\rightarrow$ FC(4) [5]

The first layer captures local spatial patterns with 5×5 kernels. Subsequent 3×3 kernels enable hierarchical feature extraction. Max-pooling reduces spatial dimensions while preserving features. Final fully connected layers aggregate spatial information. Total: $\sim 320K$ parameters, $18\times$ larger than feature-based network.

3.4 Training Protocol

Both agents used Deep Q-Learning with experience replay and target networks [2, 3]. Hyperparameters were carefully selected based on established DRL best practices [10]:

- **Learning rate:** $\alpha = 0.001$ (Adam optimizer)
- **Batch size:** 64 samples per update
- **Discount factor:** $\gamma = 0.95$ (balances immediate vs. future rewards)
- **Replay buffer:** 10,000 experiences
- **Target network update:** Every 100 episodes for DQN-Feature
- **Exploration:** ϵ -greedy with linear decay from 1.0 to 0.1 over 5,000 episodes [8]

Training consisted of 5,000 episodes per random seed, with three seeds (42, 123, 456) for statistical robustness. Evaluation episodes (100 games with $\epsilon = 0$, greedy policy only) were conducted every 1,000 training episodes to assess learned performance without exploration noise. All training was performed on CPU without GPU acceleration, demonstrating feasibility for resource-constrained settings.

3.5 Human Baseline

To contextualize agent performance, we collected human gameplay data from one experienced player over 10 rounds using pygame-based interactive control with arrow keys. The human player achieved a mean score of approximately 2.5 with scores ranging from 0 to 6. This baseline establishes a realistic performance target and enables us to determine whether learned policies achieve superhuman capability. Human performance was limited by reaction time, occasional misclicks, and suboptimal long-term planning—factors that RL agents can potentially overcome through perfect execution and learned strategies.

4 Analysis

4.1 Training Dynamics

DQN-Feature demonstrated consistent, robust improvement across all three random seeds. Evaluation performance (greedy policy, $\epsilon = 0$) increased from 2.1 ± 0.3 at episode 1,000 to final scores of 21.4 (seed 42), 21.0 (seed 123), and 13.0 (seed 456) at episode 5,000. Training loss stabilized around 0.006–0.008, indicating convergence to a stable value function approximation.

The learning trajectory exhibited three distinct phases that were consistent across seeds:

Phase 1 (Episodes 0–1,500): Initial exploration phase with predominantly random actions. Mean scores remained below 0.5 as the agent explored the state-action space. The replay buffer gradually filled with

diverse experiences, but Q-value estimates were highly unstable.

Phase 2 (Episodes 1,500–3,500): Rapid improvement phase where the agent discovered basic food-seeking behaviors. Mean scores increased from 0.5 to 2–3, surpassing human baseline performance around episode 2,000. The agent learned to move toward food and avoid immediate collisions, though long-term planning remained limited.

Phase 3 (Episodes 3,500–5,000): Policy refinement phase achieving superhuman performance (> 10 mean score). The agent developed sophisticated collision-avoidance strategies and path-planning capabilities, enabling survival for extended periods and accumulation of multiple food items per episode.

The extended training budget (5,000 vs. 1,000 episodes) proved critical for achieving superhuman performance. At 1,000 episodes, the agent barely exceeded human capability, but continued training yielded $9\times$ further improvement, demonstrating the importance of adequate convergence time.

In contrast, DQN-CNN remained near-zero performance throughout all 5,000 episodes despite its deeper architecture (3 convolutional layers). Episode scores hovered around random baseline (0.05), with no discernible learning progress. Training loss decreased initially but then oscillated without meaningful convergence. This suggests fundamental incompatibility between the CNN’s local receptive fields and the global spatial reasoning required for effective Snake gameplay at this scale.

4.2 Quantitative Comparison

Table 1: Performance metrics at episode 5,000 (mean $\pm$ std over 3 seeds, 100 eval episodes).

Metric	Feature	CNN	Human	Rand.
Score	18.5 ± 4.7	0.05	2.5	0.08
Max	46	1	—	1
Survival	629 ± 143	408	~ 1000	380
vs Human	$7.4\times$	$0.02\times$	$1.0\times$	$0.03\times$

The feature agent achieves superhuman performance ($7.4\times$ human baseline), with max score 46. All seeds substantially exceed human capability. CNN shows no learning despite 5,000 episodes.

5 Findings

Our experimental results reveal several key insights:

Superhuman performance through feature engineering: DQN-Feature achieved a mean score of 18.5 ± 4.73 across three seeds, exceeding human baseline (2.5) by a factor of $7.4\times$. The maximum individual score reached 46, demonstrating that well-designed feature representations enable reinforcement learning agents to substantially surpass human capability in constrained environments. This result validates that do-

main knowledge, when properly encoded, can dramatically accelerate learning.

Training budget criticality: Performance at 1,000 episodes (2.1 ± 0.3) only marginally exceeded human baseline, but extending training to 5,000 episodes yielded $9\times$ further improvement. This highlights the importance of adequate convergence time—premature evaluation would have underestimated the agent’s true potential. The learning curve suggests that additional training beyond 5,000 episodes may yield even higher performance.

Complete CNN failure: Despite 5,000 episodes of training, improved architecture (3 convolutional layers with increasing feature maps), and proper hyperparameter tuning, DQN-CNN remained at random baseline performance (0.05 ± 0.05). This suggests fundamental incompatibility between local convolutional receptive fields and the global spatial reasoning required for Snake. The 30×30 grid appears too large for the CNN’s receptive field to capture full game state context, yet too small to benefit from hierarchical feature learning.

Sample efficiency advantage: Feature-based representations enabled remarkably sample-efficient learning. The agent required approximately 2,000 episodes ($\sim 200K$ environment steps) to match human performance and 4,000 episodes to achieve $10\times$ human capability. In contrast, CNNs showed zero learning progress after 5,000 episodes ($\sim 500K$ steps), suggesting they would require orders of magnitude more training data to achieve comparable performance.

6 Discussion and Conclusion

6.1 Answers to Research Questions

RQ1 - State Representation Impact: Our results demonstrate that state representation choice has dramatic impact on learning outcomes. Feature-based representations dramatically outperformed CNNs, achieving scores of 18.5 vs 0.05 ($370\times$ difference).

The root causes of this disparity are threefold: (1) Feature vectors directly encoded task-relevant global information (food direction, collision proximity) that CNNs must laboriously extract through hierarchical convolutions. (2) The 30×30 grid scale occupies an awkward middle ground—too large for shallow CNNs to capture global context with limited receptive fields, yet small enough that engineered features can efficiently encode all relevant information. (3) Sample efficiency: features enabled successful learning within 2,000–4,000 episodes while CNNs showed no measurable progress after 5,000 episodes, suggesting they would require $10\text{--}100\times$ more training data.

RQ2 - Human Comparison: DQN-Feature achieved convincing superhuman performance, scoring $7.4\times$ higher than human baseline with maximum individual scores reaching 46. This demonstrates that well-designed feature representations enable RL agents to surpass human capability in constrained environments when given sufficient training time (2,000+ episodes to

match humans, 4,000+ episodes to exceed by an order of magnitude).

The agent’s superiority stems from perfect execution (no misclicks), optimal reaction time, and learned long-term planning strategies that humans struggle to execute consistently. In contrast, DQN-CNN remained substantially below human performance despite 5,000 episodes, scoring only $0.02\times$ human baseline.

RQ3 - Strategy Emergence: Feature agents exhibited clear strategy evolution across three distinct phases. In the exploration phase (episodes 0–1,500), agents displayed near-random behavior with minimal food collection. During the rapid learning phase (episodes 1,500–3,500), agents developed directed food-seeking behavior and basic collision avoidance. Finally, in the mastery phase (episodes 3,500–5,000), agents exhibited sophisticated path-planning, strategic positioning to avoid trapping themselves, and extended survival enabling accumulation of many food items per episode.

CNN agents, by contrast, showed no discernible strategy evolution. Gameplay remained indistinguishable from random throughout all 5,000 episodes, with no evidence of learned food-seeking or collision avoidance behaviors.

6.2 Theoretical Contributions

Our work challenges the prevailing assumption that CNNs universally excel at spatial reasoning tasks in reinforcement learning. We demonstrate an empirical boundary condition: while CNNs dominate in high-dimensional perceptual spaces with rich visual complexity, they can be dramatically outperformed by compact, domain-informed representations in low-to-medium complexity domains where critical information can be explicitly encoded.

For practitioners, our findings suggest that feature engineering remains highly valuable for small-scale, fully-observable environments with limited training budgets. The computational and sample efficiency advantages of features (18K vs 320K parameters, 4,000 vs >100,000 estimated training episodes) make them preferable for resource-constrained settings. However, these results do not generalize to large-scale, partially-observable, or visually-complex domains where hand-crafting features becomes infeasible.

6.3 Limitations and Future Work

Our study has several limitations that present opportunities for future research:

Grid size exploration: We tested only 30×30 grids. The CNN/feature performance crossover likely depends on grid size—smaller grids may further favor features while larger grids ($100 \times 100+$) may eventually favor CNNs as local patterns become more important.

Architectural variations: Deeper CNN architectures with attention mechanisms or larger receptive fields might overcome the limitations observed here. Similarly, recurrent architectures (LSTMs, GRUs)

could provide temporal context for both representations.

Partial observability: Testing on partially-observable Snake (limited vision radius) would better assess generalization capabilities. Features might degrade more severely than CNNs when full state information is unavailable.

Human baseline robustness: Our human baseline used only one player over 10 games. A more robust baseline with multiple players (5-10 participants, 20+ games each) would provide stronger statistical grounding.

References

- [1] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., & Riedmüller, M. (2013). Playing atari with deep reinforcement learning. *arXiv preprint arXiv:1312.5602*.
- [2] Watkins, C. J., & Dayan, P. (1992). Q-learning. *Machine Learning*, 8(3-4), 279–292.
- [3] Lin, L. J. (1992). Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine learning*, 8(3), 293–321.
- [4] Van Hasselt, H., Guez, A., & Silver, D. (2016). Deep reinforcement learning with double Q-learning. *In AAAI* (Vol. 16, pp. 2094–2100).
- [5] Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. *In Advances in neural information processing systems* (pp. 1097–1105).
- [6] Sharma, S., Sharma, M., & Atreya, G. (2017). Autonomous snake game using deep Q-learning. *arXiv preprint arXiv:1712.07124*.
- [7] Silver, D., Huang, A., Maddison, C. J., et al. (2016). Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587), 484–489.
- [8] François-Lavet, V., Henderson, P., Las Casas, D., Fonteneau, R., & Ernst, D. (2018). Introduction to deep reinforcement learning. *Foundations and Trends in Machine Learning*, 11(3-4), 219–354.
- [9] Botvinick, M., Wang, J. X., Kumaran, D., Barto, A., & Leibo, J. Z. (2020). Reinforcement learning, fast and slow. *Trends in Cognitive Sciences*, 23(5), 408–422.
- [10] Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction* (2nd ed.). MIT press.




Figure 1: Training progress (left: 50-ep moving average) and evaluation performance (right, $\epsilon = 0$). DQN-Feature improves consistently; DQN-CNN remains near zero.




Figure 2: Per-seed scores (left) and aggregate comparison (right). DQN-Feature significantly outperforms CNN and exceeds human baseline.




Figure 3: Mean episode length (survival time in steps). DQN-Feature agents survive nearly 1,000 steps (episode maximum), while DQN-CNN agents terminate prematurely (≈ 408 steps), indicating failure to learn collision avoidance.